

Free Speech Online

Philosophy, Platforms, and Controversy

Brendan Shea, PhD

Rochester Community and Technical College

What Is Free Speech?

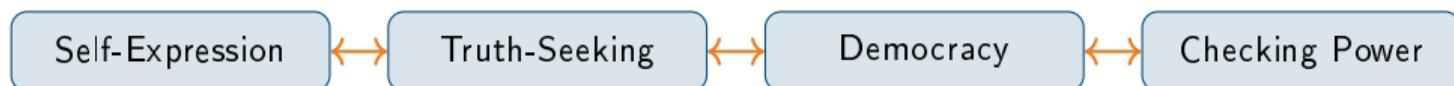
- The term **freedom of speech** refers to the right to express opinions, ideas, and information without government interference or censorship.
- Philosophers often use the broader term **freedom of expression** to include artistic expression, symbolic acts (such as burning flags or wearing armbands), and other communicative conduct beyond literal speech.
- The classic liberal conception treats free speech as a **negative liberty**, meaning freedom *from* interference by others, especially the state, rather than a positive entitlement to resources or platforms.
- Free speech protections vary enormously across legal systems, with the U.S. First Amendment providing unusually strong protections compared to most other democracies.

Key Distinction

A **negative liberty** means the state may not *prevent* you from speaking, while a **positive liberty** would mean the state must *provide* you with means to speak effectively. Most free speech traditions focus on negative liberty, but online debates increasingly raise positive questions about access to platforms and algorithmic visibility.

Why Does Free Speech Matter?

- Speaking our minds is central to who we are, so restricting speech restricts our ability to develop and express our identities—this is the value of **self-expression**.
- Open debate helps us discover truth and correct errors, since no authority is infallible enough to decide what's true for everyone—this is the value of **truth-seeking**.
- Citizens must be able to criticize government, debate policy, and persuade each other for democracy to function, making free speech essential to **democratic self-governance**.
- Free speech and a free press serve as watchdogs against government corruption, abuse, and overreach, providing a crucial mechanism for **checking power**.



Free Speech Online: What's Different?

- Online speech operates at unprecedented **scale**, with a single post capable of reaching millions of people instantly without traditional media gatekeepers.
- Information and misinformation spread with remarkable **speed**, often leaving no time for counter-speech or fact-checking before damage is done.
- Digital content has a troubling **permanence**, since posts can be archived, screenshotted, and resurface years later—the internet “never forgets.”
- The possibility of **anonymity** online enables both valuable whistleblowing and accountability-free harassment.

Feature	Traditional Media	Online Speech
Distribution	Gatekeepers (editors, publishers)	Anyone can publish instantly
Reach	Local or national audiences	Global, potentially viral
Accountability	Authors typically known	Often anonymous
Duration	Yesterday’s news	Archived indefinitely
Amplification	Editorial judgment	Algorithmic optimization

Mill's Harm Principle

- The philosopher John Stuart Mill articulated the most influential liberal framework for free speech in his 1859 work *On Liberty* (1859).
- Mill's **harm principle** holds that the only legitimate reason for society to restrict an individual's liberty is to prevent harm to others—not to protect the individual from themselves or to enforce moral conformity.
- According to Mill, the fact that others find your speech offensive, immoral, or wrong is never sufficient justification for silencing you.
- This principle places the **burden of proof** on those who would restrict speech: they must demonstrate that the speech causes genuine harm to others, not merely that it is disagreeable or unpopular.

Mill, *On Liberty* (1859)

"The only purpose for which power can be rightfully exercised over any member of a civilized community, against his will, is to prevent harm to others. His own good, either physical or moral, is not a sufficient warrant."

Mill's Marketplace of Ideas

- Mill argues that open debate is the best method for discovering truth, since no authority—whether government, church, or majority opinion—is infallible enough to decide what is true for everyone.
- Even opinions that seem obviously false should be permitted, because the process of refuting them helps us better understand *why* they are false and strengthens our grasp of the truth.
- Mill warns that beliefs held without challenge become “dead dogma”—we mouth the words but lose genuine understanding of their meaning and justification.
- This argument is often called the **marketplace of ideas**: just as economic competition produces better products, intellectual competition produces better ideas.

Mill's Fallibilism

Mill's argument rests on **fallibilism**—the recognition that any of our beliefs might be wrong. If we silence dissent, we lose the opportunity to correct our errors. Even if the silenced opinion is false, we lose the “clearer perception and livelier impression of truth, produced by its collision with error.”

Mill's Listener Autonomy

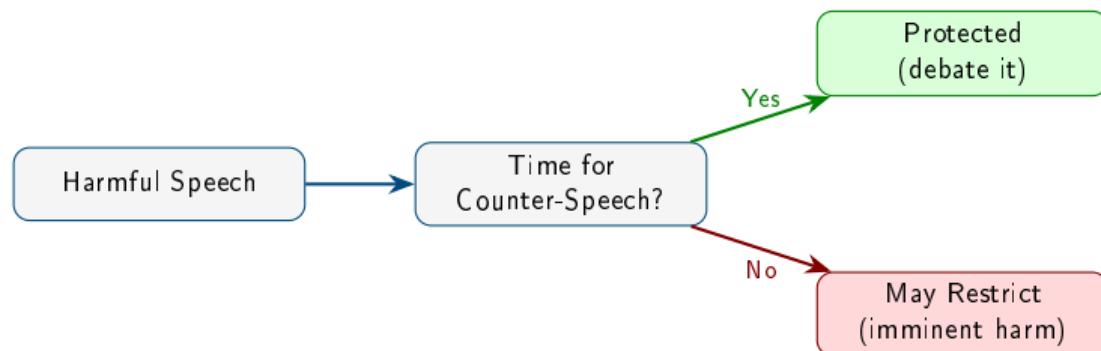
- Mill's defense of free speech emphasizes not just the rights of speakers but the rights of **listeners** to hear and evaluate ideas for themselves.
- When the state censors speech to "protect" citizens from dangerous or false ideas, it treats adults like children who cannot be trusted to think for themselves.
- Respecting **listener autonomy** means allowing people access to the full range of opinions and arguments so they can make up their own minds.
- This argument suggests that censorship is paternalistic: it assumes the government knows better than citizens what ideas they should encounter.

The Paternalism Objection

Imagine a government that bans books arguing for atheism because it believes religious faith is good for citizens. Even if the government is right that faith is beneficial, Mill would argue this censorship wrongs citizens by denying them the opportunity to consider the arguments and decide for themselves.

Applying Mill: When Is Restriction Justified?

- Mill does not believe free speech is absolute—speech that directly causes harm to others can be restricted under the harm principle.
- Mill's famous example involves a corn dealer: writing that corn dealers starve the poor is protected speech, but shouting the same words to an angry mob outside a corn dealer's house may be punished.
- The key distinction is whether there is time for **counter-speech**—if the harm is imminent and there's no opportunity for debate, restriction may be justified.
- Mill thus establishes a high bar: restriction requires not just potential harm, but harm that is direct, imminent, and unavoidable through further discussion.



The Rawls/Dworkin Approach: A Different Foundation

- Philosophers John Rawls (1971; 1993) and Ronald Dworkin (1996) offer a different justification for free speech that focuses on **equal dignity and respect** rather than on the good consequences of open debate.
- While Mill's argument is **consequentialist**—free speech is valuable because it produces good outcomes like truth and progress—the Rawls/Dworkin approach is **rights-based**.
- On this view, free speech is not just useful; it is *required* by the basic moral principle that all persons deserve equal respect as autonomous agents capable of forming their own views.

Two Different Questions

Mill asks: “Does free speech produce good results?” Rawls and Dworkin ask: “What does treating citizens as free and equal require?” Both answers support free speech, but for fundamentally different reasons.

Dworkin's Moral Independence

- Ronald Dworkin argues that individuals have a right to **moral independence**—the right not to be disadvantaged simply because others disapprove of their opinions or way of life.
- According to Dworkin, the fact that a majority finds your views offensive, immoral, or worthless is never a legitimate reason for the state to silence you.
- This is because allowing majorities to suppress minority views based on moral disapproval violates the fundamental principle that all citizens must be treated with **equal concern and respect**.

The Hard Cases

Dworkin's view is most controversial precisely where it matters most: protecting speech that most people find abhorrent. The test of a free speech principle, he argues, is whether it protects speech we hate, not speech we like.

The Legitimacy Argument

- Rawls and Dworkin also offer a **legitimacy argument** for free speech that connects it directly to the foundations of democratic government.
- Laws are only legitimate—only morally binding on citizens—if those citizens had a fair opportunity to speak against them before they were enacted.
- If the government can silence critics before a law is passed, then those who opposed the law cannot be said to have consented to it, even implicitly.

Dworkin on Legitimacy

Dworkin argues that we cannot legitimately punish someone for violating a law if they were forbidden from arguing against that law before it was passed. Free speech is what makes democratic lawmaking genuinely democratic rather than mere majority tyranny.

Discussion Questions: Philosophical Foundations

- ① Mill argues that even false and harmful opinions should generally be permitted because refuting them strengthens our understanding of truth. Do you find this convincing? Can you think of cases where exposure to false ideas makes people *less* able to recognize truth?
- ② Dworkin claims that protecting speech we find abhorrent is the real test of commitment to free speech. Is there any speech so harmful that this principle should not apply? Where would you draw the line?
- ③ The legitimacy argument says laws are only binding if citizens could speak against them. Does this apply to social media platform rules, which users had no voice in creating? Should it?
- ④ Mill wrote before the internet existed. Does the unprecedented scale and speed of online speech change his arguments? Is there still “time for counter-speech” when misinformation goes viral in minutes?

What Is Hate Speech?

- **Hate speech** refers to expression that attacks, demeans, or incites violence against people based on characteristics such as race, ethnicity, religion, gender, sexual orientation, or disability.
- Hate speech exists on a spectrum from crude slurs and stereotypes to sophisticated ideological arguments for the inferiority or dangerousness of certain groups.
- The term is contested: some argue it is too vague to regulate fairly, while others insist that the harms it causes are clear enough to justify legal restrictions.

The Spectrum of Hate Speech

Not all hate speech is equally harmful or equally easy to identify. Explicit calls for violence ("Kill all X") differ from dehumanizing rhetoric ("X are animals"), which differs from coded language and "dog whistles" that signal hateful views to insiders while maintaining plausible deniability.

The Case for Restricting Hate Speech

- Defenders of hate speech restrictions argue that such speech causes serious **dignitary harm** by denying the equal moral standing of targeted groups.
- Philosopher Jeremy Waldron (2012) argues that hate speech functions as **group defamation**—undermining the public assurance of equal citizenship.
- Perhaps most seriously, hate speech may lead to **physical violence** by normalizing hatred and dehumanizing targets—a pattern visible in genocides from the Holocaust to Rwanda.

Type of Harm	Mechanism	Example
Dignitary	Denies equal standing	Slurs, dehumanization
Psychological	Trauma, anxiety, depression	Sustained harassment
Silencing	Intimidates targets into silence	Threats, doxxing
Physical	Incites or justifies violence	Genocide propaganda

The Case Against Restricting Hate Speech

- Critics of hate speech laws argue that the concept is too **vague and subjective** to regulate fairly—what counts as “hateful” depends heavily on who is making the judgment and their political perspective.
- History shows that speech restrictions are often **used against the marginalized** they are meant to protect: early hate speech prosecutions in the U.S. targeted civil rights activists, and blasphemy laws worldwide are used to persecute religious minorities.
- Following Mill, some argue that **counter-speech is more effective** than censorship—driving hateful views underground may make them harder to monitor and refute, while public exposure allows them to be challenged.

The Slippery Slope Concern

Who decides what counts as hate speech? Today's protected minority may become tomorrow's majority with the power to silence critics. Many free speech advocates worry that hate speech laws inevitably expand beyond their original targets.

Case Study: Charlottesville and Online Radicalization (2017)

- In August 2017, white supremacists organized the “Unite the Right” rally in Charlottesville, Virginia, where counter-protester Heather Heyer was killed when a participant drove his car into a crowd.
- The rally was organized primarily through online platforms including Discord, Reddit, and 4chan, which allowed geographically dispersed extremists to coordinate and recruit new members.
- Investigators found that the attacker, James Alex Fields Jr., had been radicalized through online exposure to white supremacist content, following a pattern researchers call the **radicalization pipeline**.

Key Question

Should platforms be held responsible for hosting content that contributes to real-world violence? Or does holding them responsible create dangerous incentives to over-censor?

Case Study: Christchurch and Livestreamed Violence (2019)

- On March 15, 2019, a gunman attacked two mosques in Christchurch, New Zealand, killing 51 people—while livestreaming the massacre on Facebook Live to an audience that initially numbered in the dozens but quickly grew.
- Facebook reported that in the first 24 hours, users attempted to upload the video 1.5 million times, with 1.2 million blocked at upload—but 300,000 copies still got through and spread across platforms.
- The attack prompted the **Christchurch Call**, an international agreement between governments and tech companies to eliminate terrorist and violent extremist content online, raising questions about global internet governance.

The Virality Problem

Traditional media would never broadcast a mass shooting in real time. But social media's design—optimized for engagement and sharing—can turn atrocities into viral content before human moderators can respond.

Deplatforming: The Case of Alex Jones (2018)

- Alex Jones, host of the conspiracy theory website InfoWars, was banned from Facebook, YouTube, Apple, and Spotify in August 2018 for repeatedly violating policies against hate speech and harassment.
- Jones had promoted false claims that the 2012 Sandy Hook school shooting was a hoax and that grieving parents were “crisis actors”—claims that led his followers to harass and threaten the families of murdered children.
- Research after his deplatforming found that mentions of Jones and InfoWars across social media **dropped by 50%**, and his ability to reach new audiences was dramatically reduced.

Does Deplatforming Work?

Studies suggest deplatforming reduces reach and recruitment for extremists. But it doesn't directly change minds—Jones retained his core audience on alternative platforms. The debate is whether reducing amplification is enough.

International Approaches to Hate Speech

- The United States has the world's strongest legal protections for hate speech, with the First Amendment generally protecting even explicitly racist and hateful expression unless it constitutes a direct incitement to imminent violence.
- Most European democracies take a different approach: Germany's **NetzDG law** (2017) requires platforms to remove "manifestly unlawful" hate speech within 24 hours or face fines up to 50 million euros.
- The European Union's **Digital Services Act** (2022) creates continent-wide obligations for platforms to address illegal content, including hate speech, with significant penalties for non-compliance.
- This creates a patchwork: content legal in the U.S. may be illegal in Germany, forcing global platforms to make difficult decisions about whether to apply the strictest standard everywhere or customize content by region.

Country/Region	Hate Speech Laws?	Holocaust Denial Illegal?
United States	Limited	No
Germany	Yes	Yes
France	Yes	Yes
United Kingdom	Yes	No

The Amplification Problem

- Social media platforms don't just host speech—they actively **amplify** it through algorithmic recommendations, trending topics, and engagement-based ranking that prioritizes content likely to generate strong reactions.
- Research has consistently found that inflammatory, outrageous, and divisive content generates more engagement than measured, nuanced content—meaning platform algorithms systematically boost the most extreme voices.
- A 2021 internal Facebook study (leaked by whistleblower Frances Haugen) found that the platform's algorithms were “chroming exploiting the human brain's attraction to divisiveness” and that changes to increase engagement had made the platform angrier.

Beyond Neutrality

Platforms often claim to be neutral hosts of user content. But algorithmic amplification means they are actively making editorial choices about what speech to promote—even if those choices are made by software rather than humans.

Stochastic Terrorism

- **Stochastic terrorism** refers to using mass media to incite random acts of violence that are statistically predictable but individually unpredictable.
- The speaker doesn't direct specific attacks but creates conditions where attacks become likely, while maintaining plausible deniability.
- Critics argue the concept is too broad; defenders point to attacks from Pittsburgh (2018) to El Paso (2019) to Buffalo (2022).



$$0.001\% \times \\ 10M = 100$$

Discussion Questions: Hate Speech Online

- ① The U.S. protects hate speech that would be illegal in Germany or France. Which approach do you find more persuasive? Does the internet—which crosses all borders—require a global standard?
- ② Deplatforming Alex Jones reduced his reach by 50%. Is that a success (less exposure to harmful content) or a failure (he still has an audience and now claims persecution)? What would “success” look like?
- ③ If platform algorithms amplify hate speech because it generates engagement, are platforms morally responsible for the content they promote? Should they be legally liable?
- ④ Consider the concept of stochastic terrorism. At what point does inflammatory rhetoric become morally equivalent to incitement, even without explicit calls for violence? How would you draw that line?

The Traditional Debate Over Pornography

- For most of the 20th century, debates over pornography pitted **conservative moralists** against **liberal defenders of free expression**, with conservatives arguing that sexually explicit material corrupts morals and undermines social values.
- Liberals responded with arguments drawn from Mill: adults have the right to consume whatever content they choose, and the state should not impose majoritarian moral views on private behavior that harms no one else.
- U.S. law developed the **obscenity doctrine**, which holds that “obscene” material is not protected speech—but defining obscenity proved notoriously difficult, leading Justice Potter Stewart to famously write, “I know it when I see it.”
- The internet transformed this debate entirely: pornography went from something purchased in specialized shops to content available instantly, freely, and anonymously to anyone with a smartphone—including children.

MacKinnon's Feminist Critique

- Legal scholar **Catharine MacKinnon** (1987; 1989; 1993) reframed the pornography debate by arguing that the real harm is not moral corruption but **the subordination of women**.
- MacKinnon argues that pornography functions as a **speech act**: it doesn't merely express ideas about women but actually *does* something—it defines them as objects existing for male pleasure.
- On this view, pornography **silences women** by shaping how their speech is heard: when a woman says “no,” pornography has trained viewers to interpret refusal as part of a sexual script.
- This challenges the liberal framework directly: if pornography itself causes harm and silencing, then restricting it doesn't limit free speech—it *enables* speech.

“Pornography, in the feminist view, is a form of forced sex... an institution of gender inequality.”

—Catharine MacKinnon

Case Study: Revenge Porn and Non-Consensual Intimate Images

- **Revenge porn**—more accurately called **non-consensual intimate images (NCII)**—refers to the distribution of sexually explicit images without the subject's consent, often by former partners seeking to humiliate or control victims.
- The harms are severe and well-documented: victims report job loss, relationship destruction, depression, anxiety, and suicide attempts; the images often cannot be fully removed once they spread online.
- As of 2024, 48 U.S. states have laws against NCII, though enforcement remains difficult; the federal SHIELD Act (proposed but not yet passed) would create nationwide criminal penalties.
- Dedicated websites that host NCII have faced legal action, but content often migrates to mainstream platforms, file-sharing services, and international sites beyond U.S. jurisdiction.

Consent Is Not Transferable

A person may consent to being photographed intimately, or even to sharing images with a partner—but that consent does not extend to public distribution. The violation lies in the distribution, not the creation, of the images.

Case Study: Deepfakes and AI-Generated Intimate Images

- **Deepfakes** are AI-generated synthetic media that can place anyone's face onto explicit content, creating realistic fake pornography of people who never consented to or participated in such content.
- In January 2024, AI-generated explicit images of Taylor Swift went viral on X (formerly Twitter), viewed tens of millions of times before the platform could remove them—illustrating how quickly such content can spread.
- While celebrities have resources to fight back, ordinary people—especially women and girls—are increasingly targeted; a 2023 study found that 96% of deepfake videos online are non-consensual pornography, and the vast majority target women.

The Democratization of Harm

Creating convincing deepfakes once required significant technical skill. Today, free apps and websites allow anyone to generate explicit fake images in minutes. The barrier to creating this content has essentially disappeared.

Children and Online Pornography

- Unlike previous generations, today's children encounter online pornography not through deliberate seeking but through **accidental exposure**—pop-up ads, mislabeled content, social media links, and peer sharing on school devices.
- Research suggests the average age of first exposure to online pornography is between 11 and 13, with many children encountering explicit content even younger; most have seen pornography before receiving any formal sex education.
- Concerns focus on pornography's potential to shape developing attitudes about sex, consent, and relationships in negative ways.
- Age verification systems remain controversial: effective verification requires collecting personal data (raising privacy concerns), while ineffective systems are easily bypassed by tech-savvy children.

A Different Harm Calculus

Mill's harm principle assumes autonomous adults capable of evaluating information for themselves. Children's developing brains and lack of context change the ethical calculation—most agree that protecting children justifies restrictions that would be inappropriate for adults.

Consent in the Digital Age

- The internet has complicated **consent** in ways that earlier free speech frameworks never anticipated: content can be copied, shared, and preserved indefinitely without the subject's ongoing agreement.
- A person may consent to creating intimate content at one point in their life, but that content can resurface decades later in entirely different contexts—can consent given at 19 bind someone at 40?
- The question of **informed consent** becomes especially fraught when young people may not fully understand how permanent digital content is, how widely it may spread, or how it may affect future employment, relationships, and opportunities.

The Right to Be Forgotten

Some European courts have recognized a “right to be forgotten”—the ability to have past information delisted from search results. Should this right extend to sexual content? Can someone withdraw consent years later?

Discussion Questions: Pornography and Sexual Content Online

- ① MacKinnon argues that pornography silences women by changing how their speech is interpreted. Do you find this argument convincing? How would you test whether it's true?
- ② Should AI-generated deepfake pornography be treated the same as "real" non-consensual intimate images, even though no actual sexual act occurred? What is the nature of the harm?
- ③ OnlyFans (and related sites) creators choose to participate and control their own content. Is this genuine empowerment, or does economic pressure undermine meaningful consent? Where do you draw the line?
- ④ Most agree children should be protected from pornography. But age verification requires collecting personal data. How do we balance protecting children against privacy concerns for adults?

What Is Cancel Culture?

- **Cancel culture** refers to withdrawing support from public figures after objectionable speech or behavior—often through organized social media campaigns.
- Defenders argue cancellation is simply **accountability**—marginalized groups using collective action to impose consequences on the powerful.
- Critics argue it creates a **climate of fear**, with punishments often disproportionate and lacking due process.
- While the term was originally associated with centrist and right-wing worries about progressive activism, both historically and currently, cancellation have targeted progressives at least as often as conservatives.



Where does legitimate accountability end?

Case Study: Justine Sacco (2013)

- In December 2013, Justine Sacco, a PR executive with 170 Twitter followers, posted a tweet before boarding a flight to South Africa: “Going to Africa. Hope I don’t get AIDS. Just kidding. I’m white!”
- While she was in the air and offline for 11 hours, the tweet went viral; by the time she landed, she was the number one trending topic worldwide, had received thousands of death threats, and had been fired from her job.
- Sacco later explained the tweet was intended as **satirical commentary** on white privilege and the bubble of American ignorance about AIDS in Africa—but context collapsed in the viral spread, and the tweet was read as straightforward racism.

Disproportionate Consequences

Sacco was not a public figure, had minimal social media presence, and arguably intended the opposite of what she was accused of. Yet she faced consequences typically reserved for serious wrongdoing—job loss, death threats, lasting reputational damage.

Case Study: J.K. Rowling and the Trans Rights Debate

- Since 2019, *Harry Potter* author J.K. Rowling has made statements questioning aspects of transgender rights advocacy, including concerns about trans women in women's spaces and the medicalization of gender-questioning youth.
- LGBTQ+ advocates and many former fans have accused Rowling of **transphobia**, organized boycotts of Harry Potter products, and pressured actors from the film adaptations to distance themselves from her views.
- Rowling and her supporters argue she is raising **legitimate feminist concerns** about sex-based rights and that the backlash proves her point about the silencing of gender-critical views.

The Limits of Cancellation

Unlike Justine Sacco, Rowling has not been “cancelled” in any material sense—she remains wealthy, published, and widely read—raising questions about whether powerful people can truly be cancelled or merely criticized.

Case Study: The Assassination of Charlie Kirk (2025)

- On September 10, 2025, Charlie Kirk—the 31-year-old founder of Turning Point USA and prominent conservative activist—was shot and killed while debating students at Utah Valley University as part of his “American Comeback Tour.”
- The suspected shooter, 22-year-old Tyler Robinson, surrendered the following day; investigators found evidence suggesting **online radicalization**, with Robinson having been active in political communities on social media platforms.
- Kirk’s assassination followed a wave of political violence in 2024-2025, including two assassination attempts against Donald Trump and the killing of two Minnesota legislators.

The Pathway to Violence

FBI behavioral analysts identified a familiar pattern: grievance, fixation, validation in online communities, planning, and finally the decision that violence would solve a personal problem. Anonymous online spaces provided belonging and reinforcement.

Case Study: Academic Firings After Kirk's Death (2025)

- In the days following Kirk's assassination, dozens of professors, teachers, and university staff members were fired or placed on leave for social media posts that criticized Kirk, expressed lack of sympathy, or appeared to celebrate his death.
- Examples included a Middle Tennessee State University dean fired for posting "Looks like ol' Charlie spoke his fate into existence. Hate begets hate. ZERO sympathy," and an Iowa teacher terminated for writing "1 Nazi down."
- Republican officials actively promoted the firings: Education Secretary Linda McMahon called for more terminations, Vice President JD Vance urged citizens to report critics to employers, and Texas Governor Greg Abbott announced nearly 300 teachers were under investigation.
- Several fired educators filed federal lawsuits arguing their terminations violated the **First Amendment**, since public employees have constitutional protections for speech on matters of public concern made outside their official duties.

Platforms as the “New Governors”

- Legal scholar Kate Klonick (2018) describes major platforms as “**the new governors**”—private companies that now make more decisions about permissible speech than any government.
- Unlike governments, platforms are not bound by the First Amendment; they can restrict speech for any reason, without due process.
- This raises a fundamental tension: we rely on private companies to host public discourse, yet they have no democratic accountability.

	Governments	Platforms
Bound by First Amendment?	Yes	No
Due process required?	Yes	No
Democratically accountable?	Yes	No
Can set speech rules?	Limited	Unlimited
Scale of impact?	National	Global

Section 230: The Law That Built the Internet

- **Section 230** of the Communications Decency Act (1996) states that platforms are not legally responsible for content posted by users—they are treated as distributors, not publishers.
- This “safe harbor” allowed platforms to grow without facing lawsuits for every defamatory, harmful, or illegal post made by their billions of users.
- Section 230 also protects platforms’ right to moderate content without becoming liable for everything they miss—enabling them to remove some harmful content without assuming responsibility for all of it.
- Critics from both left and right now argue Section 230 should be reformed: conservatives claim platforms censor their views unfairly, while progressives argue platforms profit from hosting harmful content.

The 26 Words That Created the Internet

“No provider or user of an interactive computer service shall be treated as the publisher or speaker of any information provided by another information content provider.”

Discussion Questions: Cancel Culture and Platform Power

- ① Was the response to Amy Cooper proportionate justice or mob punishment? Should there be limits on how severely social media can damage someone's life for a single incident?
- ② After Charlie Kirk's assassination, conservatives used doxxing and employer pressure to get critics fired. Is this the same as "cancel culture" they previously condemned, or is celebrating an assassination categorically different?
- ③ Should platforms like Facebook and X be treated as common carriers (like phone companies) required to carry all legal speech? Or should they retain the right to set their own rules?
- ④ If Section 230 were repealed, platforms might either remove far more content (to avoid liability) or stop moderating entirely. Which outcome seems more likely? Which would be worse?

The Facebook Oversight Board

- In 2020, Facebook created an independent **Oversight Board**—sometimes called a “Supreme Court for Facebook”—to review the platform’s most controversial content moderation decisions.
- The Board includes lawyers, journalists, and human rights experts from around the world who can overturn Facebook’s decisions and issue policy recommendations.
- Its most high-profile case was upholding Facebook’s suspension of Donald Trump after January 6th, while criticizing the “indefinite” nature of the ban as lacking clear standards.
- Critics question whether a corporate-funded body can truly be independent, and whether outsourcing decisions to an unelected board solves the underlying legitimacy problem.

An Experiment in Platform Governance

The Oversight Board represents one attempt to add legitimacy and due process to platform decisions. But it reviews only a tiny fraction of cases—Facebook removes millions of posts daily that no board could possibly review.

Twitter/X Under Elon Musk

- When Elon Musk acquired Twitter in October 2022, he promised to restore it as a “free speech” platform, reinstating banned accounts including Donald Trump and various far-right figures.
- Musk dramatically reduced content moderation staff, leading to documented increases in hate speech, misinformation, and harassment on the platform (now renamed X).
- Yet Musk also suspended journalists who criticized him, banned accounts tracking his private jet, and removed content critical of governments like Turkey and India—suggesting “free speech” applied selectively.
- The X experiment illustrates how a single owner’s ideology can reshape speech rules for hundreds of millions of users overnight.

Free Speech for Whom?

Musk’s X shows that “free speech” platform policies often reflect owner preferences rather than neutral principles. One person’s free speech paradise is another’s harassment nightmare.

The Scale Problem

- Facebook users upload 350 million photos *per day*; YouTube receives 500 hours of video *per minute*; no human review system can evaluate content at this scale.
- Platforms rely on **automated moderation**—AI systems that flag or remove content based on pattern recognition, often without human review.
- These systems make inevitable mistakes:
 - **False positives** (over-moderation): Legitimate speech removed
 - News reports about violence flagged as violent content
 - Historical documentation removed as hate speech
 - **False negatives** (under-moderation): Harmful content remains
 - Novel slurs and coded language evade detection
 - Non-English content moderated less effectively
- Even 99% accuracy means millions of mistakes daily.

Who Should Decide?

- If platform executives shouldn't control speech, who should? **Governments** can regulate, but risk political censorship and vary wildly across democracies and authoritarian states.
- **Users** could vote on rules, but majorities might silence minorities—exactly what free speech principles are meant to prevent.
- **Independent bodies** like the Oversight Board add legitimacy but lack democratic accountability and can only review a tiny fraction of decisions.
- Perhaps no single answer exists: different speech contexts (news, entertainment, private messaging) may require different governance structures.

The Trilemma

We seem to want speech governance that is (1) consistent and scalable, (2) democratically legitimate, and (3) protective of minority rights. But achieving all three simultaneously may be impossible.

Key Tensions in Online Free Speech

- **Scale vs. Context:** Online speech reaches millions instantly, but meaningful evaluation of speech requires understanding context that algorithms cannot grasp.
- **Global Reach vs. Local Norms:** The internet crosses all borders, but speech norms vary dramatically—what's protected in the U.S. is illegal in Germany.
- **Private Power vs. Public Interest:** Platforms are private companies with property rights, but they control infrastructure essential to democratic discourse.
- **Protection vs. Silencing:** Restricting harmful speech protects targets, but restrictions can also silence the marginalized groups they aim to protect.

No Easy Answers

These tensions cannot be fully resolved—they must be navigated. Different societies, platforms, and contexts will balance them differently, and reasonable people will disagree.

Regulatory Approaches: How Should Platforms Be Governed?

Some proposed models for platform governance include:

- **Self-Regulation.** Allow platforms to set their own rules and enforce them as they see fit—but can profit-driven companies be trusted to balance speech and safety?
- **Government Regulation.** Require platforms to remove illegal content (EU's Digital Services Act) or verify user identity (proposed age verification laws)—but who decides what's illegal, and what about global variation?
- **Common Carrier Model.** Treat platforms like phone companies that must carry all legal speech without discrimination—but should companies be forced to amplify content they find abhorrent?
- **Public Utility Model.** Treat major platforms as essential infrastructure subject to democratic oversight—but would government control create new censorship risks?
- **Decentralization.** Replace corporate platforms with federated or blockchain-based alternatives where users control their own data—but would this make harmful content harder to moderate?

Individual Moral Obligations: What Do We Owe Each Other?

- Do we have a **duty to speak out** against injustice and misinformation, even when it makes us targets? Or does self-preservation sometimes justify silence?
- When someone is being “cancelled” for past behavior, do we have obligations to **consider context and proportionality**? Or does focusing on the accused distract from the harm they caused?
- Are we morally responsible for **how our speech is amplified** by algorithms and how others use it? If you know your criticism will trigger a harassment campaign, should you stay silent?
- Do we owe it to **future versions of ourselves** to be careful about permanent digital records? What about obligations to our past selves—can we demand old content be forgotten?
- Is there a moral obligation to **engage charitably** with those we disagree with, or does this demand unfairly burden marginalized people asked to debate their own humanity?

Conclusion: Free Speech in the Digital Age

- Mill's **harm principle** and **marketplace of ideas** remain valuable, but the internet's scale and speed challenge assumptions about counter-speech and listener autonomy.
- The **Rawls/Dworkin** framework reminds us that free speech isn't just useful—it's required by equal dignity and democratic legitimacy.
- **MacKinnon's critique** shows that some speech may itself cause harm and silence others, complicating simple free speech absolutism.
- The cases we examined—from Charlottesville to Charlie Kirk, from revenge porn to cancel culture—show that online speech raises genuinely hard questions with no perfect answers.

The Ongoing Conversation

Free speech online isn't a problem to be solved once; it's a challenge to be navigated continuously as technology evolves, norms shift, and new harms emerge. Your generation will shape what comes next.

References |

-  *Communications Decency Act of 1996, Section 230* (1996). 47 U.S.C. § 230.
-  Dworkin, Ronald (1996). *Freedom's Law: The Moral Reading of the American Constitution*. Cambridge, MA: Harvard University Press. ISBN: 978-0674319288.
-  Klonick, Kate (2018). "The New Governors: The People, Rules, and Processes Governing Online Speech". In: *Harvard Law Review* 131.6. Foundational article on platforms as private governors of online speech, pp. 1598–1670. URL:
<https://harvardlawreview.org/print/vol-131/the-new-governors-the-people-rules-and-processes-governing-online-speech/>.
-  MacKinnon, Catharine A. (1987). *Feminism Unmodified: Discourses on Life and Law*. Cambridge, MA: Harvard University Press. ISBN: 978-0674298743.
-  — (1989). *Toward a Feminist Theory of the State*. Cambridge, MA: Harvard University Press. ISBN: 978-0674896468.
-  — (1993). *Only Words*. Cambridge, MA: Harvard University Press.
-  Mill, John Stuart (1859). *On Liberty*. London: John W. Parker and Son.

References II

-  Rawls, John (1971). *A Theory of Justice*. Cambridge, MA: Harvard University Press.
-  — (1993). *Political Liberalism*. New York: Columbia University Press.
-  Waldron, Jeremy (2012). *The Harm in Hate Speech*. The Oliver Wendell Holmes Lectures. Cambridge, MA: Harvard University Press. ISBN: 978-0674065895.