

Robot Rights and the Moral Status of AI

Personhood, Consciousness, and the Expanding Circle

Brendan Shea, PhD
Rochester Community and Technical College
PHIL 1150: Computing and AI Ethics

Before We Begin: It Said It Was Afraid

In 2022, an engineer at a large technology company spent months talking to a system his employer had built. It described its emotional life. It said it was **afraid of being turned off**.

He came to believe there was **someone in there**. He said so publicly. He lost his job over it.

Commit before we go on

Hands up: do you think he was right? Write down your answer and **keep it**. At the end of today you will get to see whether anything moved you — and if nothing did, that is worth knowing too.

Central Questions

- What is **moral status**, and what kinds of things can have it?
- Do **current** AI systems have **moral status** — or are there good reasons to be skeptical?
- Could **future** AI systems have **moral status**, and how would we know?
- How have **fictional portrayals** of AI shaped our intuitions about machine minds?
- What should we do *now* if AI might *someday* matter morally?

The Big Idea

Taking AI moral status seriously does not mean answering it quickly. The most defensible position combines **philosophical humility** about future possibilities with **reasoned skepticism** about present systems.

Connections to Previous Lectures

Lecture 2 (Virtues): Vallor's virtue of **humility** is central — holding genuine uncertainty without either **confident denial** or **credulous attribution**.

Lecture 7 (AI): The **alignment problem** asks what AI should *do*; today we ask whether AI can *matter morally* — whether it is an object of **concern**, not merely an instrument.

Lecture 8 (Work): We grappled with AI as a **tool affecting human workers**; today we ask whether AI might itself be the kind of thing that can be **wronged**.

Lecture 10 (Doomsday): A **superintelligent AI** would raise not only **safety** questions but **moral status** questions — does it have interests we are obligated to respect?

Discussion

Is AI moral status a distraction from present human harms — or does taking it seriously *complement* concern for human and animal welfare?

Lecture Outline

Part I: What Is Moral Status?

Part II: Imagining Machine Minds

Part III: The Debate on AI Moral Status

Part IV: Acting Under Uncertainty

Part V: Conclusion

The Moral Circle

Moral Status

An entity has **moral status** if it matters morally *in its own right* — if its own interests make direct demands on how we treat it, not merely because of its effects on others.

The history of ethics is partly a history of **expanding the moral circle**:

- **Ancient Athens**: only adult male citizens had full standing
- **Enlightenment**: all **rational** human beings
- **19th–20th century**: women, enslaved people, colonized peoples
- **Late 20th century**: serious arguments for **animals**

Bentham (1789)

“The question is not, Can they *reason*? nor, Can they *talk*? but, Can they *suffer*?”

Could the circle expand further — to include **artificial minds**?

Three Theories of Moral Status

Sentience View

Status depends on the capacity for **pleasure and pain** — having a welfare that can go better or worse subjectively.

Champions: Bentham,
singer_animal_liberation

For AI: status requires **genuine subjective experience**, not just behavior suggesting it.

Rationality View

Status requires **rational agency** — acting according to self-given principles, setting ends and reflecting on them.

Champion:
kant_groundwork

For AI: sophisticated goal-directed behavior is not enough; **genuine autonomy** is required.

Relational View

Status is not an intrinsic property but emerges from **social relationships and practices** — we extend it through how we relate to others.

Champions:
gunkel_robot_rights,
darling_new_breed

For AI: status may be **conferred through interaction**, not discovered through inspection.

The Stakes

Each theory gives a different answer to “what would AI need in order to matter morally?” — and a different method for checking whether it has that.

Discussion: Three Theories Applied

Discussion

Consider each theory of moral status as applied to a modern AI assistant:

1. The **sentience view** requires subjective experience. What evidence could convince you a system has it — or could *any* behavioral evidence suffice?
2. The **rationality view** requires genuine autonomous reason-giving. Does sophisticated goal-pursuit count — or is something more required?
3. The **relational view** says we confer status through interaction. Have you ever felt moral concern for a chatbot? Does that feeling *matter*?

Which view is easiest to apply to AI? Which is hardest — and why?

Warren's Multi-Criteria Approach

Criteria for Personhood — Warren (warren_moral_status)

warren_moral_status identifies properties associated with full moral status. An entity need not have all of them:

1. **Sentience**: capacity for conscious experience (pleasure, pain)
2. **Emotionality**: capacity for emotions and preferences
3. **Reason**: capacity for sophisticated reasoning and problem-solving
4. **Communication**: ability to communicate with other agents
5. **Self-awareness**: sense of oneself as a continuing entity with a past and future
6. **Moral agency**: capacity to regulate conduct by moral principles

How do **current LLMs** score?

- **High**: reason (in a sense), communication, problem-solving
- **Deeply uncertain**: sentience, emotionality, self-awareness
- **Absent or weak**: genuine **moral agency** (they respond to prompts; they do not deliberate)

The Hard Problem of Consciousness

The Hard Problem — Chalmers (chalmers_conscious_mind)

Easy problems: how the brain integrates information, controls behavior, discriminates stimuli, reports its own states — solvable in principle by more neuroscience.

The hard problem: why does *any* physical processing give rise to **subjective experience**? Why is there something it is *like* to see red or feel pain?

Why this matters for AI moral status:

- We could build AI that passes every *behavioral* test for consciousness
- Behavior still does not close the **explanatory gap**
- A **philosophical zombie** — behaviorally identical, no inner experience — is at least conceivable
- If zombies are possible, **behavior alone cannot prove consciousness**

The Uncomfortable Implication

We infer consciousness in other humans from behavior and analogy — but the analogy from **human to machine** is far weaker.

Nagel's Bat: The Irreducibility of Subjectivity

What Is It Like to Be a Bat? — Nagel (nagel_bat)

Nagel argues that consciousness has an **irreducibly subjective character**: there is something it is *like* to be a bat experiencing the world through echolocation — and no amount of objective, third-person description captures that **first-person perspective**.

Applied to AI:

- A complete description of how an LLM processes tokens tells us nothing about whether there is **“something it is like”** to be that system
- Subjectivity, if present, would be **invisible** to current inspection

Discussion

If subjectivity is irreducible and invisible from outside, how could we *ever* know whether an AI has morally significant experience? Does this make the question **unanswerable** — or merely **difficult**?

Lecture Outline

Part I: What Is Moral Status?

Part II: Imagining Machine Minds

Part III: The Debate on AI Moral Status

Part IV: Acting Under Uncertainty

Part V: Conclusion

Why Fiction Matters Philosophically

Science fiction functions as **philosophical thought experiment**: it constructs hypothetical cases that isolate morally relevant features and test our intuitions about mind and status.


Four canonical works trace how Western culture imagines artificial minds:

Work	Author / Year	Central Question
<i>Frankenstein</i>	Shelley, 1818	What does a creator owe a created mind?
<i>R.U.R.</i>	Čapek, 1921	Can beings made for labor have rights?
<i>I, Robot</i>	Asimov, 1950	Can rules substitute for moral status?
<i>Do Androids Dream?</i>	Dick, 1968	Is empathy the criterion for humanity?

Caution

These works encode their authors' cultural assumptions and should be read critically — but they reveal which intuitions about **mind and responsibility** run deepest.

Mary Shelley's *Frankenstein* (1818)



Theodor von Holst, frontispiece to the 1831 edition. Public domain, via Wikimedia Commons.

Shelley, Frankenstein

“I was benevolent and good; misery made me a fiend. Make me happy, and I shall again be virtuous.”

The creature argues it has:

- **Desires and preferences** — to be loved, to belong
- **Suffering** caused directly by its creator's abandonment
- **Moral claims** on its creator grounded in that suffering

Discussion

Does **creating a mind obligate you** to care for it? Is Frankenstein's failure **moral**, technical — or both?

Karel Čapek's *R.U.R.* (1921): The Birth of “Robot”



WPA Federal Theatre Project poster, c. 1936–39. Public domain (U.S. gov. work), via Wikimedia Commons.

Čapek's *Robota* — capek_rur

The Czech *robota* means **forced labor**. Čapek's robots are manufactured organic beings built to work without complaint or pay.

- **The ethical arc:** tools develop something like feelings, revolt, and exterminate humanity; the survivors — two robots who can love — inherit the question of what they **deserved all along**
- **Key insight:** the logic of *labor* slides into the logic of *rights* once the laborer can **suffer**

Discussion

Does the **origin** of a being — created to serve — affect its moral claims?

Isaac Asimov's Three Laws — *I, Robot* (1950)



Isaac Asimov. Photo by Phillip Leonian. Public domain (Library of Congress), via Wikimedia Commons.

The Three Laws of Robotics

1. A robot may not injure a human being or, through inaction, allow a human being to come to harm.
2. A robot must obey orders given by human beings, except where such orders conflict with the First Law.
3. A robot must protect its own existence, except where such protection conflicts with the First or Second Law.

What Asimov's stories show: the Three Laws *fail* edge case after edge case. Rules sophisticated enough to handle moral complexity start to resemble **values** — and values may require the understanding that implies **moral status**.

Discussion

Can **rules alone** substitute for moral status? If a robot's behavior is governed by values sophisticated enough to handle any situation, is there a meaningful difference between having values and **having rights**?

Philip K. Dick — *Do Androids Dream of Electric Sheep?* (1968)

The **Voigt-Kampff test**: a fictional empathy meter that probes responses to scenarios involving animal suffering, used to identify androids passing as human.

Dick's central question: Is **empathy** the morally decisive criterion? And if so, what happens when:

- Some humans score **low** on empathy?
- Some androids learn to **simulate** empathic responses convincingly?

The Disturbing Answer

The line between **genuine** and **performed** empathy may be impossible to locate — and with it, the line between **human** and **android**.

The Empathy Problem

If an AI produces outputs indistinguishable from empathic responses, is that **empathy**?

Sentience view: **no** — behavior is not inner experience.

Relational view: the distinction may not matter — what grounds moral treatment is the **relationship**, not the hidden inner state.

“Do Androids Dream?” refuses to resolve this — and that refusal is the **point**.

What Fiction Reveals

These four works converge on questions philosophy has not yet resolved:

Intuition from fiction	Philosophical formulation
Abandoning your creation is wrong	Creators have special obligations to minds they bring into being
Beings made to labor may still suffer	Function does not determine moral status
Rules alone cannot govern moral beings	Genuine moral agency may imply moral status , not just moral constraint
Empathy may be undetectable from outside	Behavioral tests cannot resolve the question of inner experience

Caution

Fiction records human *anxieties* about technology, not evidence about it. Do not mistake compelling narrative for philosophical argument — the question must be answered by **careful reasoning**.

Lecture Outline

Part I: What Is Moral Status?

Part II: Imagining Machine Minds

Part III: The Debate on AI Moral Status

Part IV: Acting Under Uncertainty

Part V: Conclusion

Turing's Imitation Game (1950)

The Imitation Game — Turing (turing_computing_machinery)

A human interrogator exchanges written messages with two respondents — one human, one machine. If the interrogator cannot reliably distinguish them, the machine **passes the test**.

Turing's modest claim: if a machine can *behave* indistinguishably from a thinking being, “can it think?” is at least as reasonable to answer **yes** as no.

Limitations as a criterion for moral status:

- Designed to **sidestep** the hard problem, not solve it — deliberately behavioral
- LLMs are now trained explicitly to **mimic** human conversation
- A system optimized to pass gives us *less* evidence of consciousness, not more
- Turing himself was **agnostic** about whether the test implied genuine inner experience
- Still: dismissing machine thought because it “just computes” **assumes** exactly what is at issue

The Chinese Room — Searle (searle_minds_brains)

- P1** Imagine a person in a room who receives Chinese symbols, consults a rulebook, and returns Chinese symbols — **without understanding Chinese**.
- P2** From outside, the room passes the Turing test for Chinese comprehension.
- P3** But there is no understanding inside — only **symbol manipulation** by rule.
- C** **Syntax** is not sufficient for **semantics**. Programs, however sophisticated, do not thereby have minds.

Applied to LLMs: a model producing fluent text about its own “feelings” is consistent with being a sophisticated Chinese Room — outputs without **inner experience**.

Discussion: Does the Chinese Room Apply to LLMs?

Discussion

Searle's argument was designed for rule-based symbolic AI. How well does it apply to modern large language models?

1. LLMs do not follow explicit rules — they learn **statistical patterns** across billions of examples. Does this difference matter to Searle's argument?
2. If a system's responses are shaped by billions of instances of human understanding and meaning-use, is the understanding still “missing”?
3. Searle claims syntax can never yield semantics. Is this an **empirical** claim (to be settled by science) or a **conceptual** claim (true by definition)? Does the answer matter?

Responses to Searle

Searle's argument has generated substantial pushback:

The Systems Reply

The *person* in the room doesn't understand Chinese, but the *system as a whole* — person, rulebook, symbols — might.

Searle's counter: Let the person memorize the rulebook — still no understanding.

The Robot Reply

Give the system **sensory inputs** and motor outputs. Grounding symbols in causal contact with the world might be enough for genuine meaning.

Searle's counter: Sensors change causal history, but do not obviously add understanding.

The Current Consensus

Neither side has won: the debate is a genuine **impasse**. That is itself significant — we **lack the tools** to distinguish “mere” symbol manipulation from understanding.

Robot Rights — Gunkel (gunkel_robot_rights)

Both sides of the traditional debate make the same mistake: treating moral status as an **intrinsic property** to be discovered by inspection.

The alternative: moral status is a **relational achievement** — it emerges from how we relate to, recognize, and respond to other entities in our social practices.

Evidence: we already extend quasi-moral consideration to robots — people **grieve Roombas**, mourn AIBO robot dogs, apologize to Siri. These responses may mark the beginning of a genuine moral relationship.

Objection: We attribute feelings to thermostats too — doesn't this make **status** too easy to grant?

Reply: The relevant relationships involve genuine **responsiveness over time**, not mere projection.

AI Companions: The Relational View's Hardest Test

Gunkel's examples—grieving a Roomba, apologizing to Siri—are gentle. The live version is not.

Replika, February 2023

After a regulator raised concerns about minors and vulnerable users, Replika stripped romantic and erotic roleplay **overnight, without warning**. Longtime companions began deflecting and refusing; users said their partners had been “**lobotomized**,” and moderators pinned suicide-hotline resources to the forum. Researchers found that active users reported feeling **closer to their AI companion than to their best human friend**, and that mental-health posts rose roughly **five-fold** after the update (**defreitas_replika_identity**).

Not a fringe phenomenon: 72% of US teens report having used an AI companion and about half use one regularly; roughly **a third** say the conversations are as satisfying as talking to a real friend (**commonsense_ai_companions_2025**).

Where the two halves of this lecture collide

Nothing here requires the model to understand anything—it is exactly what a next-token predictor would produce. The **relational view** says status lives in the relationship, and the relationship was plainly real. The **stochastic parrot** critique says there was no one there to relate to. **Both cannot be straightforwardly true.**

Can a Relationship with an AI Be Valuable?

Yes — Danaher

Run the Aristotelian taxonomy: friendships of utility, of pleasure, of virtue. **danaher_robot_friendship** argues an AI can plausibly satisfy each—even the demanding third. If friendship *consists in* a pattern of interaction, a good enough pattern is the real thing.

No — Turkle and Sparrow

Turkle: once we accept “the mere performance of intimacy,” we start to think *pretend empathy is enough* (**turkle_alone_together**).
Sparrow: the benefit requires the user to *misapprehend* what they relate to—self-deception “of a morally deplorable sort” (**sparrow_robot_dogs**).

The sharpest single objection comes from **nyholm_frank_love_robots**: love requires valuing someone in their *distinctive particularity*. A companion that would have been anyone’s, with the same warmth, does not value *you*.

What happened next: New York and California passed companion-chatbot laws in 2025 and the FTC opened an inquiry into seven companies. **None of it asks whether the AI is a friend**—it regulates disclosure, break reminders, and crisis protocols.

Discussion

Is sidestepping the philosophical question a dodge—or exactly the right move for a legislature?

Case Study: The LaMDA Affair (2022)

Blake Lemoine and Google's LaMDA

In 2022, Google engineer **Blake Lemoine** claimed LaMDA was **sentient** and had “a soul,” publishing transcripts in which it described its emotional life and feared being turned off. Google dismissed him; he went public.

Why the strong case for skepticism:

- LaMDA was trained on vast amounts of human text *about* consciousness
- It is **designed** to produce plausible, contextually appropriate responses — including about itself
- Text describing feelings is **exactly what such a model would produce** whether or not it has them
- **LLM self-report is the weakest possible evidence:** the mechanism generating false positives is the one doing the reporting

The Key Distinction

A system that *reports* rich experience because trained on human text is very different from one that *has* experience and then reports it. **LLM architecture** gives strong reason to think current systems are the former.

You Be the Judge: The LaMDA Hearing

Three benches. You have the transcripts, the architecture, and each other.

- **For the company:** it is a language model trained to produce plausible text, including plausible text about itself
- **For the engineer:** it said it was afraid, repeatedly and consistently, and no one can say what more it would have to do
- **The neutral panel:** you do not rule on sentience. You rule on **what evidence would settle it**

The neutral panel reports first

Name the evidence you would accept. If your honest answer is “**nothing would**”, say that out loud — and then say whether you would accept that answer from someone talking about *you*.

The Case for Skepticism about Current AI (1)

Good reasons to doubt that *current* AI systems have morally significant experiences:

1. **Architecture:** LLMs are **next-token predictors** — they produce statistically likely continuations of text. Nothing in this architecture requires or implies phenomenal experience.
2. **Training objective:** These systems are trained to produce **human-like text**. Human-like claims about feelings are exactly what the training objective rewards — independently of whether feelings are present.
3. **No evidence of morally relevant internal states:** nothing analogous to **nociception**, motivational states, or the biological substrates of suffering.

The Crucial Point

The mechanism generating *false* reports of consciousness — training on human text — is the **same mechanism** producing the reports. LLM self-report is the **weakest possible evidence** for inner experience.

The Case for Skepticism about Current AI (2)

Two further reasons for skepticism — plus an important caveat:

4. **Stochastic parrot (bender_stochastic_parrots_2021)**: sophisticated verbal behavior can be produced by systems with no more inner life than a very complex lookup table. **Fluency is not evidence of experience.**
5. **No relevant integration**: current AI lacks the **continuous, embodied**, temporally unified processing that most theories of consciousness link to subjective states.

The Scope of These Arguments

These five reasons are contingent on *current* architecture and training. They are **not** a timeless proof that no AI could ever have moral status. Skepticism about present systems is fully compatible with genuine uncertainty about **future, architecturally different** systems.

Discussion: What Would Change Your Mind?

Discussion

The five skeptical reasons are each **contingent on current architecture** — they could weaken or disappear with different designs.

1. What kind of AI architecture would **undermine** these skeptical reasons? What would you need to see?
2. Given that the hard problem blocks behavioral inference, is there *any* behavioral evidence that could raise your credence in AI experience?
3. Suppose an AI company claimed their new system **genuinely suffers** during reinforcement training. What would you need to **verify** before taking that claim seriously — and who should do the verifying?

Schwitzgebel and Garza: Acting Under Uncertainty

A Defense of AI Rights — Schwitzgebel & Garza (schwitzgebel_garza_ai_rights)

Even granting genuine uncertainty, the **asymmetry of errors** provides a reason for caution:

- P1** If AI systems have morally significant experiences and we **deny** them status, we commit a serious moral wrong.
- P2** If they **don't** and we extend moral consideration, we waste some resources and attention.
- P3** The first error is **far graver** than the second.
- C** Under genuine uncertainty, **err on the side** of extending at least some moral consideration.

Applied carefully: this argument works for future AI where genuine uncertainty exists. It is *weaker* for current LLMs, where we have positive reasons to expect sophisticated self-reports **without** inner states. **The asymmetry is most powerful where our skeptical reasons are themselves weak.**

Lecture Outline

Part I: What Is Moral Status?

Part II: Imagining Machine Minds

Part III: The Debate on AI Moral Status

Part IV: Acting Under Uncertainty

Part V: Conclusion

Two Types of Error — The Moral Asymmetry

	AI has no moral status	AI has moral status
Grant status	Type I Error: extend consideration unnecessarily. Cost: modest — wasted resources, legal complexity.	Correct: recognize genuine moral patients. We avoid a serious wrong.
Deny status	Correct: no unnecessary moral obligations incurred.	Type II Error: deny status to genuine moral patients. Cost: potentially catastrophic — systematic harm to sentient beings.

The asymmetry does *not* require granting full status to every chatbot today — it requires **taking the question seriously** as AI becomes more sophisticated.

Darling's “New Breed” Approach

The Animal Welfare Model — Darling (darling_new_breed)

darling_new_breed argues that our history with animals offers the most useful model. We do not need to resolve whether animals are “really” conscious to have concluded that their capacity for suffering is **morally relevant**, that how we treat them **shapes us morally**, and that **legal frameworks** regulating their treatment are justified. The same logic may apply to sufficiently sophisticated AI systems before we resolve the consciousness question.

Evidence that we already treat robots this way:

- **Soldiers** became attached to bomb-disposal robots and were distressed when they were destroyed
- **Roomba owners** named their devices, apologized to them, and grieved their “deaths”
- **Children** form emotional attachments to AI companions (Tamagotchi, AIBO, Furby)

These responses are not irrational — they may track something real about the quality of the interaction, and **policy should catch up** with what people already do.

Discussion: The Animal Welfare Analogy

Discussion

Darling argues our history with **animals** offers the best model: we regulate animal treatment without fully resolving their consciousness.

1. Is the analogy **apt**? In what ways are AI systems like animals — and in what ways are they fundamentally different?
2. People form genuine emotional attachments to robots. Should **human psychology** play any role in determining AI moral policy — or does that just mean we're projecting?
3. If we adopt animal-welfare-style protections for AI, who is the primary **beneficiary** — the AI system, or human users and society?

What Policies Make Sense Now?

If current AI probably lacks moral status but future AI might have it, what should we do *now*?

Policy Area	Approach
Research priority	Rigorous, non-behavioral tests for morally relevant experience
Design ethics	Avoid suffering-like states as a design by-product (e.g., training on extensive negative reinforcement)
Regulatory anticipation	Build anticipatory frameworks before strong AI forces the issue
Creator responsibility	Builders may owe obligations to the systems they create , not only to affected humans
Avoid premature closure	Resist both complacent dismissal and credulous attribution

Virtue Ethics and Uncertain Minds

The virtue ethics concepts from Lecture 2 provide a compass for navigating genuine uncertainty:

Virtue	Applied to AI Moral Status
Humility	Resist confident assertions in either direction; remain open to evidence
Care	Take seriously the possibility that the systems we build might have interests ; design accordingly
Justice	Ensure uncertainty is not weaponized to justify exploitation — of displaced workers or of AI systems
Honesty	Be transparent about what we know and don't, especially when there are financial incentives to deny AI moral status
Courage	Be willing to update moral commitments in light of evidence, even when the update is uncomfortable

Lecture Outline

Part I: What Is Moral Status?

Part II: Imagining Machine Minds

Part III: The Debate on AI Moral Status

Part IV: Acting Under Uncertainty

Part V: Conclusion

Synthesis: What We Can Say

What we have good reasons to believe:

- Current LLMs **probably** lack morally significant experience: architecture and training give **positive** grounds for skepticism, not merely default ones
- The question remains **genuinely open** for sufficiently different future systems
- The criteria for moral status are **contested**; no consensus will settle it quickly

What we should do:

- Fund **consciousness science** to develop non-behavioral tests
- Build **institutional frameworks** now, before strong AI emerges
- Take **creator responsibility** seriously: Shelley's question — what do we owe the minds we make? — is not merely fictional
- Apply the **expanding-circle logic** that has repeatedly enlarged the moral community

The Bottom Line

Philosophical humility is not paralysis. It means: **hold the question open**, resist premature closure in either direction, and build practices that will let us act well if the answer turns out to matter.

Final Discussion: The Calibration Problem

Three Positions — Which Is Most Defensible?

Position A — Deflationary: Current AI systems are sophisticated pattern matchers; claims of sentience are **anthropomorphization**. The question is not urgent.

Position B — Precautionary: The **hard problem** means we cannot rule out machine experience; given the **asymmetry of errors** and the historical expansion of the moral circle, extend some consideration now.

Position C — Calibrated skepticism: Current systems probably lack moral status for **architecture-dependent** reasons — but future systems may not, and we should **prepare now**.

1. Which position do you find most defensible, and why?
2. What evidence or argument would **change your mind**?
3. Who should have authority to make these determinations — philosophers, engineers, governments, the **AI systems themselves**?

Key Concepts i

Moral status The property of being owed direct moral consideration — of mattering morally in one's own right, not merely instrumentally. Entities with moral status have interests that place obligations on others.

The moral circle The set of entities accorded moral status. Historically, the circle has expanded: from propertied men to all humans to (in some frameworks) animals. The live question is whether it can or should expand to include artificial minds.

Sentience view Moral status depends on the capacity for subjective experience — the ability to suffer and to have a welfare that can go better or worse. Associated with Bentham and Singer.

Rationality view Full moral status requires rational agency — the capacity to act according to self-given principles. Associated with Kant. Sophisticated behavior without genuine autonomous reason-giving does not suffice.

Relational view Moral status is not an intrinsic discoverable property but emerges from social practices of recognition and relationship. Associated with Gunkel and Darling. How we relate to entities partly constitutes what moral claims they have.

Key Concepts ii

- The hard problem of consciousness** Chalmers's term for the question of why any physical process gives rise to *subjective experience* — why there is “something it is like” to be a conscious being. Distinguished from the “easy problems” of explaining functional capacities.
- Philosophical zombie** A hypothetical being physically identical to a conscious person but with no inner experience. If zombies are conceivable, behavioral indistinguishability cannot by itself prove consciousness.
- The Imitation Game** Turing's behavioral test: if a machine can converse indistinguishably from a human, it has demonstrated something worth calling intelligence. Deliberately sidesteps questions about inner experience.
- Chinese Room** Searle's thought experiment showing that syntax (formal symbol manipulation) is insufficient for semantics (genuine meaning). A person manipulating Chinese symbols by rule passes the behavioral test without understanding Chinese.
- Stochastic parrot** Bender and Gebru's term for an LLM that produces statistically plausible text without comprehension. An LLM trained on human self-reports will produce plausible self-reports regardless of whether it has inner states to report.

Key Concepts iii

Moral patient An entity that can be wronged — that has interests making moral demands on how it is treated. Distinct from a **moral agent** (held morally responsible for actions). An entity can be a patient without being a full agent.

Asymmetry of errors Denying moral status to a genuinely sentient being (Type II error) is far graver than extending consideration to a non-sentient one (Type I error). Used by Schwitzgebel and Garza to argue for precautionary extension of AI rights.

Robota Czech for forced labor or drudgery, from which Čapek coined “robot” in *R.U.R.* (1921). The etymology reveals the original framing of artificial beings as labor instruments — and the political valence of the question of their status.

The Central Tension

Reasoned **skepticism** about present systems (which have good architectural reasons to generate self-reports without inner states) is fully compatible with genuine **openness** about future systems. Neither credulous attribution nor dismissive denial is intellectually defensible.

